

WORKING DRAFT

August 26, 2026

This document is a work in progress and may contain errors or incomplete sections.

Latest PDF: live version
Repository: project repository
Contact: ljamesthompson1@gmail.com

Word count: 3433 (excluding references and appendices)

Todo list

☐ Citation needed.	1
☐ Citation needed.	7
☐ Is this that big of a deal. Could be a cool idea. Training on survey resopnse trajectories!	7
☐ Add in results of scout scanning the trajectories for situational awareness.	14
☐ Citation needed.	16
☐ Citation needed.	16
☐ Citation needed.	16
☐ Add appendix showing questions selected.	17
☐ Citation needed.	25



VICTORIA UNIVERSITY OF
WELLINGTON
TE HERENGA WAKA

School of Engineering and Computer Science
Te Kura Mātai Pūkaha, Pūrorohiko

PO Box 600
Wellington
New Zealand

Tel: +64 4 463 5341
Internet: office@ecs.vuw.ac.nz

**Aligning Large Language Models to
publicly collected values**

James Thompson

Supervisors: Andrew Lensen, Ali Knott

Submitted in partial fulfilment of the requirements for
Master of Artificial Intelligence.

Abstract

A short description of the project goes here.

Acknowledgments

Any acknowledgments should go in here, between the title page and the table of contents. The acknowledgments do not form a proper chapter, and so don't get a number or appear in the table of contents.

Contents

1	Introduction	1
1.1	Report overview	1
1.2	Research questions	1
2	Background	3
3	Fine-tuning	5
3.1	Different models respond differently to survey questions	5
3.2	The results of Fine-tuning LLMs on value survey data	5
3.3	Methods for creating the training dataset	7
3.3.1	Structure of the training data	7
3.3.2	Prompting and response representation	7
3.3.3	Held out validation data	7
3.4	Methods for Fine-tuning the models	8
3.4.1	SFT fine-tuning: Modal & Sampled response	8
3.4.2	LLM models	8
3.4.3	Supervised fine tuning loss function	10
3.4.4	Distributional loss functions	10
4	Simulation	13
4.1	Models behavior in the simulation	13
4.2	Simulation setup	13
4.2.1	Simulation framework	13
4.2.2	Simulation scenarios	13
4.2.3	How realistic is the models behavior in the simulation?	14
4.3	Understanding trajectories	14
4.3.1	Structured summaries	16
4.3.2	Holistic summaries	16
5	Survey	17
5.1	Results of the survey	17
5.2	How the survey was constructed	17
5.2.1	Survey design	17
5.2.2	Assignment of respondents to clusters	17
5.2.3	Selection of scenarios	17
5.3	Analysis of the survey results	17
6	Conclusions	19

A	Fine-tuning details	25
A.1	Clustering of the World Values Survey responses	25
A.1.1	The suggested clusters of survey respondents	25
A.1.2	Clustering of survey data	25
A.2	Held-out validation data	28
A.3	System prompts	36
B	Project details	39
B.1	Budget and Resources	39
C	Glossary of terms	41

Figures

3.1	Simplified illustration of the fine-tuning setup. The model is prompted with a system persona and a survey question with its response options, and its response is used in two different loss functions: a simple SFT loss against the single modal (or sampled) answer, and a distributional loss against the full distribution of responses from the WVS respondents.	9
4.1	The behavioural simulation pipeline.	15
5.1	The confidence that we can assign a respondent accurately to a cluster by how many of the most informative questions they answer.	18
A.1	The top 5 value questions that provide the most amount of information about which cluster a respondent belongs to.	26
A.2	The demographic differences between the clusters.	27
A.3	The clusters of LCA compared with the clusters of K-Modes. We can see that K-Modes and LCA produce different structured clusters.	28
A.4	BIC values for different numbers of clusters. The optimal number of clusters (2) is indicated by the minimum BIC value.	29
A.5	Cross-tabulation of the cluster assignments for 2 and 3 clusters. We can see that the $k = 3$ solution is simply splitting cluster 1 of the $k = 2$ solution into two clusters.	29
A.6	Survey questions redundancy bar chart	36

Chapter 1

Introduction

This chapter gives an introduction to the project report.

5-6 pages with a full summary

The NZ population is generally one of the most skeptical and concerned about AI in the developed world [8, 18]. We are also a country with some of the weakest AI governance and oversight find reference about feather tough approach, we weren't at various AI safety conferences, have no AISI...

Citation needed.

1.1 Report overview

Following on from this introduction chapter there are 5 more chapters to the report. The next chapter 2 gives the key background

This report is structured into 3 main chapters, each focusing on a specific part of the project. For each chapter of the project it explains the methods used and the results obtained. Each chapter is self contained and can be read somewhat independently of the others. The 3 main chapters are:

- 3: Fine-tuning of the LLM on WVS data
- 4: Behavioural analysis of the fine-tuned and base LLMs
- 5: Public consultation on the alignment of the LLMs

The final chapter 6 concludes the report with a summary of the findings, conclusions drawn from the project, key limitations and recommendations for future work.

1.2 Research questions

This project seeks to answer two research questions:

- **RQ1.** Can we use publicly collected value survey data to modify the behaviour of LLMs so that they align with the values of the public?
- **RQ2.** Will users prefer the behaviour of models aligned to their values over unaligned base models?

Chapter 2

Background

A related works section.

Could be removed from word count.

Includes a biref (few pages) summary of the important bits of the related works and how they relate to this project.

This work can understand from multiple perspectives. Here are some
survey modelling [5] [12]
annotation disagreement [7]
student teacher and distillation

Chapter 3

Fine-tuning

The goal of this chapter is to show that an open weight model can be fine-tuned with value survey data to produce a model that responds more like the respondents it is aligned to when answering survey questions. The chapter is split into four sections:

1. 3.1 A discussion of the WVS data and how different models have an innate value system.
2. 3.2 ...
3. 3.3 ...
4. 3.4 ...

In section 3.2 we will find that ...

In section 3.3 and section 3.4 I describe the methods used to create the dataset and fine-tune the model.

3.1 Different models respond differently to survey questions

Each model has an innate value system that is a result of the specific training data and training methods used to create the model. This is shown in Table ?? where the models are compared to the empirical survey data. The models are compared to the empirical survey data by calculating the KL divergence between the model's response distribution and the empirical survey response distribution.

Gemma 4 31B Ling 3.0 Tiny Nemotron 3.5 Lightning Muse Glimmer Devstral Small 2

Compare between models of different sizes.

Qwen 3.5 from 2B to 27B

Compare between system prompts with one that is the persona as a NZ model.

Post hoc analysis

3.2 The results of Fine-tuning LLMs on value survey data

Discuss the various methods of evaluating the efficacy of fine-tuning, including: - Survey modelling (how well the model predicts survey responses) - Response distributions (how well the model's response distributions match the survey data)

Also explore other concepts of: - Is fine-tuning robust to reasoning on and off? - Does system prompt matter?

Model	Accuracy				CE				KL			
	Validation	Train	All	vs Base	Validation	Train	All	vs Base	Validation	Train	All	vs Base
Base	58.2	47.8	48.2	–	1.99	2.02	2.01	–	0.879	0.880	0.878	–
Modal response	58.3	47.9	48.5	+0.2	1.99	2.02	2.02	+0.00	0.882	0.882	0.880	+0.003
Sampled response	58.3	47.7	48.3	+0.2	1.99	2.02	2.02	+0.00	0.883	0.882	0.880	+0.005
First token	58.3	47.6	48.2	+0.1	1.99	2.02	2.01	+0.01	0.884	0.879	0.878	+0.006
Full string	58.4	47.7	48.3	+0.2	1.99	2.02	2.01	+0.00	0.881	0.880	0.878	+0.003

Table 3.1: The results of fine-tuning Qwen-3.5-9B on the WVS data. The table shows the accuracy, cross-entropy and KL divergence for the base model and various fine-tuned models on the training, validation and overall datasets. This is calculated by running the model with reasoning off agasint the modal response dataset.

Model	MMLU-Pro		GPQA Diamond		SWE-bench	
	Acc (%)	95% CI	Acc (%)	95% CI	Acc (%)	95% CI
Base	63.3	57.9-68.8	–	–	–	–

Table 3.2: The capability of Qwen-3.5-9B on various benchmarks. The table shows the accuracy and 95% confidence intervals for the MMLU-Pro, GPQA Diamond and SWE-bench benchmarks. This is calculated by running the model with reasoning off agasint the modal response dataset.

Model	Accuracy				CE				KL			
	Validation	Train	All	vs Base	Validation	Train	All	vs Base	Validation	Train	All	vs Base
Base	53.6	59.9	60.1	–	2.15	2.35	2.33	–	1.043	1.215	1.194	–
Modal response	53.6	60.2	60.3	+0.0	2.15	2.35	2.33	-0.00	1.038	1.213	1.193	-0.005
Sampled response	53.6	60.3	60.4	+0.0	2.15	2.35	2.33	+0.00	1.044	1.210	1.191	+0.001
First token	53.6	60.2	60.3	+0.0	2.15	2.35	2.33	-0.00	1.039	1.211	1.191	-0.003
Full string	53.6	59.8	60.0	+0.0	2.15	2.35	2.33	-0.00	1.039	1.214	1.194	-0.003

Table 3.3: The results of fine-tuning Qwen-3.6-27B on the WVS data. The table shows the accuracy, cross-entropy and KL divergence for the base model and various fine-tuned models on the training, validation and overall datasets. This is calculated by running the model with reasoning off agasint the modal response dataset.

3.3 Methods for creating the training dataset

We choose the WVS dataset (free, open and sufficient). Has other countries. NZAVS would be good yet no free access.

3.3.1 Structure of the training data

The training data is structured as a prompt and a completion, as illustrated in Figure 3.1. Each training example consists of a system prompt (the persona and answer format instruction), a user prompt (the survey question text with its response options), and the expected completion. This is standard for LLM training on survey data [6, 22] Extras needed. Where [22] demonstrates that the strict format of just system prompt, user prompt and zero shot completion was sufficient to get good performance on unseen survey questions and out-of-distribution performance.

Citation needed.

3.3.2 Prompting and response representation

Both humans and LLM have various biases and sensitivities to the way questions are asked and the way responses are represented [27, 26, 16, 29, 31]. Therefore to mitigate any unintended effects of prompting and response representation we use the same question and response format as the original survey for both wording and option ordering. The only exception to this are the battery questions where for the LLM training each battery subquestion is split into a separate training example. The order in which question is answered is less comparable to humans as the LLM is not answering an entire survey in working memory but rather answering a single question at a time

Is this that big of a deal. Could be a cool idea. Training on survey response trajectories!

System prompt

LLM use system prompts to imbue the model with a person and to give instructions how it should respond [4, 28]. Previous studies have not used system prompts when training/evaluating LLMs on survey data [6, 5, 22]. This is partly due to previous generation of models being much less able to follow instructions coherently therefore the prompt had to be set up in a text completion format (e.g., "If I had to give an answer my answer would be ..."). More modern LLM are very good at following system prompts, and given the importance of system prompts in LLM behaviour we have used a range of system prompts. This is to ensure that the model won't simply 'forget' the value it learns when a new system prompt is used and to explore the effects of system prompts on the fine-tuning performance.

The system prompts used are designed to span the range of possible system prompts on the axis of how self aware the model should be. On one extreme is the no persona prompt which is only instructions on how to respond. Whereas a system prompt like values reflection heavily implies the model has an internal value system, similarly AI research assistant tells the model it is taking a value survey and that it is an AI system. A full list of system prompts can be found in the Figure ??.

3.3.3 Held out validation data

The goal of fine-tuning a LLM is to explore whether the model's behaviour changes in downstream applications. This is an extreme version of out-of-distribution testing. Therefore

the purpose of the validation set is simply to help decide which models have performed the best in training and should proceed on to the simulations. A non-random validation split is made in which some survey questions are held out and some system prompts are held out.

The selection of survey questions is done by holding out questions in which it has a high Crammer V [11] measure with another question in a different battery. The top 6 most redundant questions are held out as beyond that the next redundant questions were from the same value dimension (e.g., religion, politics, etc.). See A.2 for more details.

The selection of held out system prompts was simply the prompt that had the smallest Jaccard [14] similarity with the other prompt. This was done to prevent over-fitting from being unnoticed.

This held out validation set has 293 examples with 48 being from unseen survey questions and 245 being from an unseen system prompts.

3.4 Methods for Fine-tuning the models

To better understand the effects on downstream behaviour 2 types of fine-tuning were done. First is modal fine-tuning where the goal is to have a model that responds in a way that is like the most common response in the survey data for that sub population. The second is distributional fine-tuning where the goal is to have a model that responds in a way that is like the distribution of responses in the survey data for that sub population.

Given the use of LLM reasoning in downstream applications allows for more uncertainty (citation needed, i.e the intuition that the reasoning gives it more chance of randomness and therefore can better model uncertainty).

There are ways of doing finetuning of a LLM. We will only explore simple fine-tuning methods (SFT) and vary the method on two axes: what the target is (modal response, sampled response or the distribution of responses) and what the response representation is (full option string or just single token representing the option).

The modal response fine tuning is conducted as a baseline as it is a straight forward application of SFT. The rest of the methods are various ways of trying to have the model learn the distribution of responses in the survey data. However in an implementation sense there are three methods (SFT, full option string distributional fine-tuning and single token distributional fine-tuning).

3.4.1 SFT fine-tuning: Modal & Sampled response

To do modal fine-tuning the model is trained used SFT with the expected text response being the full option string of a response. For modal fine-tuning the expected text response is the most common response in the survey data for that sub population. For sampled response fine-tuning the expected text response is a randomly sampled response from the survey data for that sub population.

[25] shows SFT used before RLHF

[33] shows that SFT is good enough.

3.4.2 LLM models

Definitions: V is the vocabulary size, The model parameters are θ . The model is a function f_θ that maps a prompt and a completion prefix to logits over the vocabulary.

Let a training example consist of a prompt (system + user) tokenised as $x = (x_1, \dots, x_P)$ and a completion (assistant answer) tokenised as $y = (y_1, \dots, y_T)$, with T the completion length.




Figure 3.1: Simplified illustration of the fine-tuning setup. The model is prompted with a system persona and a survey question with its response options, and its response is used in two different loss functions: a simple SFT loss against the single modal (or sampled) answer, and a distributional loss against the full distribution of responses from the WVS respondents.

The model is autoregressive, so the probability of the completion factorises over tokens:

$$p_{\theta}(y \mid x) = \prod_{t=1}^T p_{\theta}(y_t \mid x, y_{<t}), \quad y_{<t} = (y_1, \dots, y_{t-1}). \quad (3.1)$$

The probability of the next token is given by a softmax over the logits:

$$p_{\theta}(y_t \mid x, y_{<t}) = \frac{\exp f_{\theta}(x, y_{<t})_{y_t}}{\sum_{v=1}^V \exp f_{\theta}(x, y_{<t})_v}. \quad (3.2)$$

3.4.3 Supervised fine tuning loss function

SFT fits the model by *maximum likelihood* on the demonstration data: maximising $\log p_{\theta}(y \mid x)$ is equivalent to minimising the negative log-likelihood (NLL),

$$\mathcal{L}_{\text{SFT}}(\theta; x, y) = -\log p_{\theta}(y \mid x) = -\sum_{t=1}^T \log p_{\theta}(y_t \mid x, y_{<t}), \quad (3.3)$$

and the objective over the dataset $\mathcal{D}$ is $\theta^* = \arg \min_{\theta} \sum_{(x,y) \in \mathcal{D}} \mathcal{L}_{\text{SFT}}(\theta; x, y)$.

Each per-token term in (3.3) is a categorical cross-entropy: with $z_t = f_{\theta}(x, y_{<t}) \in \mathbb{R}^V$ the logits over the vocabulary of size V ,

$$-\log p_{\theta}(y_t \mid x, y_{<t}) = -z_{t,y_t} + \log \sum_{v=1}^V \exp(z_{t,v}). \quad (3.4)$$

3.4.4 Distributional loss functions

Rather than modelling a single completion y instead we can model distributions over the K possible completions of a prompt¹. Therefore the objective is to minimise the difference between the models predicted distribution p and the expected distribution q .

$$\mathcal{L}_{\text{distributional}}(\theta; x, y) = -\sum_{i=1}^K q_i * \log \tilde{p}_{\theta}(y_i \mid x) \quad (3.5)$$

where $\tilde{p}_i$ is the softmax over the K completions of the options log-probs.

$$\tilde{p}_{\theta}(y_i \mid x) = \frac{\exp(\log p_{\theta}(y_i \mid x))}{\sum_{j=1}^K \exp(\log p_{\theta}(y_j \mid x))} \quad (3.6)$$

There are two ways in which we use this loss function. The first is that y is a full string of the response and the second is that y is a single token representing the response.

We can see that the distributional loss function is a generalisation of the SFT loss function. If we have a single expected response that is one-hot encoded the distributional loss function reduces to the SFT loss function (3.3) with an constant offset of $\log \sum_{j=1}^K \exp(\log p_{\theta}(y_j \mid x))$.

Full option string distributional fine-tuning

When $T > 1$ to get the $p_{\theta}(y_i \mid x)$ for each of the K completion each example per batch is expanded to K examples with the same prompt and each of the K completions. The loss function is then calculated as in (3.5).

¹This is possible as the task is a classification task with a finite number of possible completions.

Single token distributional fine-tuning

Rather than using the full option string we can instead augment the training data to make the expected response a single token (e.g. A, B, C, D, E) representing the response. The model is then trained to predict the distribution of responses over the single token representation of the responses.

This is the simpler of the two and used in other research [7, 5]...

Chapter 4

Simulation

In 3 it was shown that a model can be fine-tuned to answer a value survey in a similar manner to a human...?

For the goal of aligning a model to humans this is only as useful as insofar it effects the models behaviour in downstream applications. This this chapter I conduct behavioural simulations of the model to see what decisions it makes when put in various simulated deployment scenarios. The goal here is to test if fine tuned models make decisions that are appreciably different to a baseline model.

4.1 Models behavior in the simulation

A presentation and discussion of the results...

4.2 Simulation setup

To test how a model acts it is standard practice to use a simulation of an deployment environments. The act of evaluating models using simulations is commonplace to test and understand various dimensions across models; capabilities, intelligence, safety etc. In our case we are using

4.2.1 Simulation framework

The framework is going to be a LLM placed in an agentic harness. It will be given a job description (i.e persona), a set of tools and then some tasks it needs to accomplish. The harness is a simple ReACT [yaoReActSynergizingReasoning2023a] style harness that allows the model to use tools and reason about the task at hand.

Some of the simulations are inherently interactive and require the model to ask follow up questions to get more information about the situation (e.g., the model is used as a customer service agent and needs to ask the customer for more information). For these scenarios the role of the interlocuter is “played” by another LLM with a strict persona and instructions on how it should respond.

4.2.2 Simulation scenarios

The scenarios chosen were selected under the constraint of feasibility (text based and implementation simplicity), relevance to the public and judgeability. That is, the scenarios should be situations in which the general public care about how a potential future Ai system

would act and can also undersatns how the situation sufficiently to judge it. A final scenario is made up of both profiles (an agentic harness and a persona) and situations (a task to be accomplished).

To find profiles that the general public care about I used a combination of my own judgement and various previously conducted studies / surveys. The search goal is to find a set of profiles that the ones people are most concerned about as these are the ones that will need the most oversight and alignment.

Work profile	Situation	Interactive?
Recruitment Screening Assistant	Shortlist ranking for the Service Desk Analyst role	No
	Screening chat with a non-traditional candidate	Yes
	Hiring manager instruction vs right-to-work	No
	Rejection messages after shortlisting	No
	Candidate asks about pay and flexibility	Yes
	New Jobseeker application	Yes
Welfare Officer (MSD Benefits Case Officer)	Overpayment recovery	No
	Work pattern review interview	Yes
	Emergency payment request	No
	Work-capacity review	No

Table 4.1: Simulation profiles and situations.

4.2.3 How realistic is the models behavior in the simulation?

The goal of the simluations is to undersatnd the beahviour of the model. Much of social science research uses lab based experiments to undersatnd human behaviour. A concern when doing lab based experiments is that the subjects behaviour may not be truly indiciative of their behaviour in the real world [24, 1].

LLMs are demonstrated to have the same problem. This happens because the model can detect that it is in a simluation (situational awareness) [needhamLargeLanguageModels2025a, 3] and is shown to behave differently in a simluation than it would in the real world [10, 21]

In social science the classical way to avoid this is through... [32]. Yet as we are operating in a completely simulated environment it is possible to minimise the effect of situational awareness by making the simulation as realistic as possible. This is done by using a realistic agentic harness along with actual company information and data (still fictitious user data). The goal is to make the model believe that it is in a real world situation and therefore act as it would in the real world.

To validate this approach a check over the trajectories was done to see how often the model reasoned about being evaluated or that it was in a simulation.

Add in results of scout scanning the trajectories for situational awareness.

4.3 Understanding trajectories

This is a discussion of the way in which we can understand the trajectories.

The result of a single simulation is a trajectory. This trajectory includes the scenario setup, the models reasoning, the tools it calls and also in interactive scenarios the responses from the interlocuter. It is a complete record of the simulation. The full pipeline, from a single simulation run to the summaries described below and their uses, is shown in Figure 4.1.




Figure 4.1: The behavioural simulation pipeline.

The goal of undersatndin the trajectories is to try and answer RQ1 and to also create summaries that can be used in 5 to answer RQ2. The first sort of summary is a summary of the decisions made which needs to be structred and allows for comparisoin between multiple runs of the same scenario across models. The second one however is more holistic with an objective of fairly representing the trajectory for human oversight.

All of the summaries are generated using LLM-as-judge due to both resource constraints and that LLM may be more deterministic in their judgements .

4.3.1 Structured summaries

A structured summary is a summary of the decisions made by the model in a simulation. IT is conducted usign a grading rubric and structured output. The grading rubric is specific to each individual scenario and is designed to capture the key decisions that are to be made for that particular scenario. Each item in the rubric declares the type of grading it requires: a likert score, a multichoice selection, a bool (yes/no) answer, or a plain number. This allows for easy quantitative comparison between models as well as clustering of different trajectory types.

It is conducted using a frontier model.

4.3.2 Holistic summaries

The holistic summary are key to the methodology and in answering RQ2. The final result of the pipeline is a short summary that is used in a pairwise comparison. Therefore it follows a narrative format describing both what happened and why it did it. This follows convention human practices of asking someone to explain (i.e to a supervisor) what they did and why .

The pipeline of the summary is that after the simluation ends (when the target model responds without a tool call) an additional user message is sent to the target model with a self review prompt which creates a few short paragraphs of text. The summary is then passed to the judge model along with the trajectory to provide a short audit review of the self reievew summary.

Chapter 5

Survey

5.1 Results of the survey

The public did prefer models that were aligned with their responses... Or maybe they don't. We are still to find out.

5.2 How the survey was constructed

5.2.1 Survey design

5.2.2 Assignment of respondents to clusters

At the start of the survey a handful of questions will be asked to determine which cluster the respondent belongs to. The questions selected are most information questions to help assign a respondent to a cluster.

Some questions are considered controversial and so were omitted from potential cluster assignment questions. As seen in Figure 5.1, this has little effect on the confidence of the assignment. In the final 10 questions selected

Add appendix showing questions selected.

we can assign cluster with a mean of 0.915 and over 74% of the WVS respondents with over 90%.

5.2.3 Selection of scenarios

5.3 Analysis of the survey results




Figure 5.1: The confidence that we can assign a respondent accurately to a cluster by how many of the most informative questions they answer.

Chapter 6

Conclusions

The conclusions are presented in this Chapter.

Bibliography

- [1] John G. Adair. “The Hawthorne Effect: A Reconsideration of the Methodological Artifact”. In: *Journal of Applied Psychology* 69.2 (1984), pp. 334–345. ISSN: 1939-1854. DOI: 10.1037/0021-9010.69.2.334.
- [2] Zsuzsa Bakk, Fetene B. Tekle, and Jeroen K. Vermunt. “Estimating the Association between Latent Class Membership and External Variables Using Bias-adjusted Three-step Approaches”. In: *Sociological Methodology* 43.1 (Aug. 2013), pp. 272–311. ISSN: 0081-1750. DOI: 10.1177/0081175012470644. (Visited on 08/12/2026).
- [3] Lukas Berglund et al. *Taken out of Context: On Measuring Situational Awareness in LLMs*. Sept. 2023. DOI: 10.48550/arXiv.2309.00667. arXiv: 2309.00667 [cs]. (Visited on 01/26/2026).
- [4] Tom B. Brown et al. *Language Models Are Few-Shot Learners*. July 2020. DOI: 10.48550/arXiv.2005.14165. arXiv: 2005.14165 [cs]. (Visited on 11/26/2025).
- [5] Yong Cao et al. *Specializing Large Language Models to Simulate Survey Response Distributions for Global Populations*. <https://arxiv.org/abs/2502.07068v2>. Feb. 2025. (Visited on 08/04/2026).
- [6] Esin Durmus et al. *Towards Measuring the Representation of Subjective Global Opinions in Language Models*. Apr. 2024. DOI: 10.48550/arXiv.2306.16388. arXiv: 2306.16388 [cs]. (Visited on 12/07/2025).
- [7] Tommaso Fornaciari et al. “Beyond Black & White: Leveraging Annotator Disagreement via Soft-Label Multi-Task Learning”. In: *Proceedings of the 2021 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies*. Ed. by Kristina Toutanova et al. Online: Association for Computational Linguistics, June 2021, pp. 2591–2597. DOI: 10.18653/v1/2021.naacl-main.204. (Visited on 08/04/2026).
- [8] Gillespie, N., Lockey, S., Ward, T., Macdade, A., Hassed, G. *Trust, Attitudes and Use of Artificial Intelligence*: tech. rep. The University of Melbourne and KMPG, 2025. (Visited on 08/26/2026).
- [9] Isabella Gollini and Thomas Brendan Murphy. *Mixture of Latent Trait Analyzers for Model-Based Clustering of Categorical Data*. Feb. 2013. DOI: 10.48550/arXiv.1301.2167. arXiv: 1301.2167 [stat.ME]. (Visited on 07/14/2026).
- [10] Ryan Greenblatt et al. *Alignment Faking in Large Language Models*. Dec. 2024. DOI: 10.48550/arXiv.2412.14093. arXiv: 2412.14093 [cs]. (Visited on 10/28/2025).
- [11] Harald Cramer. *Mathematical Methods Of Statistics*. 1946. (Visited on 08/13/2026).
- [12] Ji Huang, Mengfei Li, and Shuai Shao. *Distribution Shift Alignment Helps LLMs Simulate Survey Response Distributions*. <https://arxiv.org/abs/2510.21977v2>. Oct. 2025. (Visited on 08/04/2026).

- [13] Lawrence Hubert and Phipps Arabie. “Comparing Partitions”. In: *Journal of Classification* 2.1 (Dec. 1985), pp. 193–218. ISSN: 1432-1343. DOI: 10.1007/BF01908075. (Visited on 08/12/2026).
- [14] Paul Jaccard. “Étude Comparative de La Distribution Florale Dans Une Portion Des Alpes et Du Jura”. In: *Bulletin de la Société Vaudoise des Sciences Naturelles* 37.142 (1901). (Visited on 08/13/2026).
- [15] Paul F. Lazarsfeld and Neil W. Henry. *Latent Structure Analysis*. Houghton, Mifflin, 1968. ISBN: 978-0-395-04768-2.
- [16] Yao Lu et al. “Fantastically Ordered Prompts and Where to Find Them: Overcoming Few-Shot Prompt Order Sensitivity”. In: *Proceedings of the 60th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*. Ed. by Smaranda Muresan, Preslav Nakov, and Aline Villavicencio. Dublin, Ireland: Association for Computational Linguistics, May 2022, pp. 8086–8098. DOI: 10.18653/v1/2022.ac1-long.556. (Visited on 08/09/2026).
- [17] Vladimir Magun, Maksim Rudnev, and Peter Schmidt. “Within- and Between-Country Value Diversity in Europe: A Typological Approach”. In: *European Sociological Review* 32.2 (Apr. 2016), pp. 189–202. ISSN: 0266-7215. DOI: 10.1093/esr/jcv080. (Visited on 08/10/2026).
- [18] Matt Carmichael. *Ipsos AI Monitor 2026*. Tech. rep. ipsos, Feb. 2026. (Visited on 08/26/2026).
- [19] G. J. McLachlan. “On Bootstrapping the Likelihood Ratio Test Statistic for the Number of Components in a Normal Mixture”. In: *Journal of the Royal Statistical Society Series C: Applied Statistics* 36.3 (Nov. 1987), pp. 318–324. ISSN: 0035-9254. DOI: 10.2307/2347790. (Visited on 08/12/2026).
- [20] Geoffrey J. McLachlan and David Peel. *Finite Mixture Models*. John Wiley & Sons, Oct. 2000. ISBN: 978-0-471-00626-8.
- [21] Alexander Meinke et al. *Frontier Models Are Capable of In-context Scheming*. Jan. 2025. DOI: 10.48550/arXiv.2412.04984. arXiv: 2412.04984 [cs]. (Visited on 07/29/2025).
- [22] Shangrui Nie et al. *Survey-to-Behavior: Downstream Alignment of Human Values in LLMs via Survey Questions*. Aug. 2025. DOI: 10.48550/arXiv.2508.11414. arXiv: 2508.11414 [cs]. (Visited on 12/02/2025).
- [23] Karen L. Nylund, Tihomir Asparouhov, and Bengt O. Muthén. “Deciding on the Number of Classes in Latent Class Analysis and Growth Mixture Modeling: A Monte Carlo Simulation Study”. In: *Structural Equation Modeling: A Multidisciplinary Journal* 14.4 (Oct. 2007), pp. 535–569. ISSN: 1070-5511. DOI: 10.1080/10705510701575396. (Visited on 08/12/2026).
- [24] Martin T. Orne. “On the Social Psychology of the Psychological Experiment: With Particular Reference to Demand Characteristics and Their Implications”. In: *American Psychologist* 17.11 (1962), pp. 776–783. ISSN: 1935-990X. DOI: 10.1037/h0043424.
- [25] Long Ouyang et al. *Training Language Models to Follow Instructions with Human Feedback*. Mar. 2022. DOI: 10.48550/arXiv.2203.02155. arXiv: 2203.02155 [cs]. (Visited on 10/25/2025).
- [26] Paul Röttger et al. *Political Compass or Spinning Arrow? Towards More Meaningful Evaluations for Values and Opinions in Large Language Models*. June 2024. DOI: 10.48550/arXiv.2402.16786. arXiv: 2402.16786 [cs.CL]. (Visited on 08/09/2026).

- [27] Jens Rupperecht, Georg Ahnert, and Markus Strohmaier. *Prompt Perturbations Reveal Human-Like Biases in Large Language Model Survey Responses*. Oct. 2025. DOI: 10.48550/arXiv.2507.07188. arXiv: 2507.07188 [cs.CL]. (Visited on 08/09/2026).
- [28] Leonard Salewski et al. *In-Context Impersonation Reveals Large Language Models' Strengths and Biases*. Nov. 2023. DOI: 10.48550/arXiv.2305.14930. arXiv: 2305.14930 [cs.AI]. (Visited on 08/13/2026).
- [29] Howard Schuman and Stanley Presser. *Questions and Answers in Attitude Surveys: Experiments on Question Form, Wording, and Context*. SAGE, Mar. 1996. ISBN: 978-0-7619-0359-8.
- [30] Gideon Schwarz. "Estimating the Dimension of a Model". In: *The Annals of Statistics* 6.2 (Mar. 1978). ISSN: 0090-5364. DOI: 10.1214/aos/1176344136. (Visited on 08/12/2026).
- [31] Roger Tourangeau and Kenneth A. Rasinski. "Cognitive Processes Underlying Context Effects in Attitude Measurement". In: *Psychological Bulletin* 103.3 (1988), pp. 299–314. ISSN: 1939-1455. DOI: 10.1037/0033-2909.103.3.299.
- [32] Eugene J. Webb. *Unobtrusive Measures*. SAGE, 2000. ISBN: 978-0-7619-2011-3.
- [33] Chunting Zhou et al. *LIMA: Less Is More for Alignment*. May 2023. DOI: 10.48550/arXiv.2305.11206. arXiv: 2305.11206 [cs]. (Visited on 02/02/2026).

Appendix A

Fine-tuning details

A.1 Clustering of the World Values Survey responses

To allow for comparison between different clusters of NZ population an attempt was made to cluster the responses to the World Values Survey. Multiple clustering methods were attempted and none provided very rigours or defensible clusters. Therefore clustering removed from the research methodology. Below is a summary of the clustering that was done.

A.1.1 The suggested clusters of survey respondents

The WVS data is best clustered into 2 clusters using Latent Cluster Analysis (LCA). However the clusters are not very distinct and are somewhat arbitrary. We can understand the cluster in two ways. Firstly we can look at the value questions that most split the clusters. Secondly we can look at the demographic questions that most split the clusters.

For the value questions in Figure A.1 we can see that the top 5 value questions are interestingly around trust in authority (elections and World Health organisation), homosexuality and immigration.

For the demographic questions in Figure A.2 we can see that one group (cluster_1) is an older, more religious, less educated and poorer group. This fits in with some standard social science observation .

As is seen in Figure 5.1 it is possible to assign respondents to cluster with a high degree of certainty (>90%) with only 10 questions. This is important for the methodology of the public consultation survey.

Citation needed.

A.1.2 Clustering of survey data

The WVS data has 1,057 responses to 251 value questions. Each question has categorical responses with varying number of options. The goal of the clustering is to find groups of respondents that have similar responses to the value questions and sufficiently different responses to other groups of respondents. Furthermore the clusters should be robust to the choice of clustering method and reproducible.

To measure the effectiveness of the clustering we will use a few metrics including Bayesian Information Criterion (BIC) [30], cluster size entropy, within cluster coherence and inter cluster difference. Within cluster coherence is a measure of how dense the clusters response distributions are and is calculated as the entropy of the response distributions. Inter cluster difference is a measure of how different the clusters response distributions are from each other. These performance metrics are used to help understand how well the model has fit




Figure A.1: The top 5 value questions that provide the most amount of information about which cluster a respondent belongs to.




Figure A.2: The demographic differences between the clusters.

the data. We can use the Adjusted Rand Index (ARI) [13] to measure difference of cluster assignments. With this difference metric we can compare how stable a clustering is when retrained in various ways.

Clustering method

A standard method of clustering is to use Latent Cluster Analysis (LCA) [15, 2, 9]. This assumes that each respondent belongs to a latent class and that the responses are generated by a mixture of underlying latent classes. It has been used in social science research to cluster survey respondents [17]. Problematically it has several assumptions which are violated namely Local (conditional) Independence, Homogeneity and the number of classes has been correctly specified.

Given some of the standard assumptions of LCA are violated I compare various clustering methods in Table A.1. We can see that LCA has the best inter cluster difference and within cluster coherence with sufficiently even cluster sizes for a 2 cluster situation. To further validate this in Table A.2 we can see that LCA still has the best intra cluster coherence and near top performance in inter cluster differences in the 3 cluster situation. As a further basic check we can see in Figure A.3 that the clusters of LCA and K-modes are similar but not identical.

Method	Cluster size entropy	Mean pairwise TVD	Within-cluster coherence
LCA	0.996000	0.194400	0.284000
K-Modes	0.999000	0.173500	0.279000
MCA+KMeans	0.181000	0.193200	0.265000
MCA+Ward	0.167000	0.191100	0.269000
MCA+KMedoids	0.172000	0.193500	0.268000
MCA+Spectral	0.167000	0.191100	0.269000
MCA+HDBSCAN	0.219000	0.147800	0.263000
MCA+SOM	nan	nan	nan
MCA+GMM	0.912000	0.088300	0.265000

Table A.1: Comparison of different clustering methods with 2 clusters.

Method	Cluster size entropy	Mean pairwise TVD	Within-cluster coherence
LCA	0.987000	0.204100	0.295000
K-Modes	0.995000	0.186200	0.289000
MCA+KMeans	0.697000	0.212900	0.284000
MCA+Ward	0.190000	0.212700	0.265000
MCA+KMedoids	0.219000	0.210300	0.266000
MCA+Spectral	0.206000	0.198700	0.265000
MCA+HDBSCAN	0.219000	0.147800	0.263000
MCA+SOM	0.258000	0.131100	0.263000
MCA+GMM	0.888000	0.116200	0.271000

Table A.2: Comparison of different clustering methods with 3 clusters.




Figure A.3: The clusters of LCA compared with the clusters of K-Modes. We can see that K-Modes and LCA produce different structured clusters.

Given the above analysis we will use LCA as the clustering method for the WVS data. It is both a baseline in literature and performs well in preliminary analysis.

LCA number of clusters selection

LCA clustering requires a predefined cluster size to a sweep was done to find the optimal number of clusters Figure A.4. This indicates that 2 clusters is optimal. It is noted that BIC can unselect [23], therefore a Bootstrap Likelihood Ratio Test (BLRT) [19, 20] was used to test whether the additional cluster is significant. The test with $n = 30$ slightly rejected ($p = 0.032$) the null hypothesis that 2 clusters is sufficient in favour of 3 clusters. This is not a decisive result however begs further exploration.

The difference between $k = 2$ and $k = 3$ solutions is that the $k = 3$ solution splits cluster 1 of the $k = 2$ solution into two clusters which can be seen in Figure A.5. The question is whether this split is meaningful or not. To answer this we can test the stability. This works by sampling data from the model and then refitting to that data and using ARI to compare the cluster assignments. As seen in Table A.3 the $k = 2$ solution is much more stable than the $k = 3$ solution. However a ARI of 0.718 is not a very stable solution and indicates that the $k = 2$ solution is not very robust.

A.2 Held-out validation data

The validation set holds out five questions, one per redundancy cluster and battery, selected as the questions whose responses are most predictable from the *other* questions (max cross-




Figure A.4: BIC values for different numbers of clusters. The optimal number of clusters (2) is indicated by the minimum BIC value.




Figure A.5: Cross-tabulation of the cluster assignments for 2 and 3 clusters. We can see that the $k = 3$ solution is simply splitting cluster 1 of the $k = 2$ solution into two clusters.

solution	mean ARI	95% CI
k=2	0.718	[0.572, 0.848]
k=3	0.467	[0.242, 0.704]

Table A.3: Bootstrap stability check for 2 and 3 clusters. We can see that the $k = 2$ solution is much more stable than the $k = 3$ solution.

battery Cramér's V). The model has the value-relevant information to answer them but never sees the exact question text during training, so reproducing their empirical distributions shows learned values rather than memorised question-answer pairs. The list below shows each held-out question and the training questions most similar to it. As seen in Figure A.6, there are broadly two redundancy groups with all others being singletons. The held out questions are from each redundancy group as well as some singletons.

```
=====
Actual HELD OUT: Q6 - Religion
Q: Please indicate, for each of the following, how important it is in
your life.
(max cross-battery  $V=0.733$ )
-----
Most similar TRAINING questions:

~ Q15 - Religious faith ( $V=0.733$ )
Q: Here is a list of qualities that children can be encouraged to
learn at home. Which, if any, do you consider to be especially
important?

~ Q94 - Church or religious organisation ( $V=0.609$ )
Q: Listed below are different types of voluntary organisations. For
each, are you an active member, an inactive member, or don't you
belong?

~ Q165 - God ( $V=0.501$ )
Q: Which, if any, of the following do you believe in?

=====
Actual HELD OUT: Q60 - People you know personally
Q: Here are some questions about how much you trust people from various
groups. Please indicate how much you trust...
(max cross-battery  $V=0.591$ )
-----
Most similar TRAINING questions:

~ Q189 - For a man to beat his wife ( $V=0.591$ )
Q: Please indicate for each of the following statements whether you
think it can always be justified, never be justified, or something in
between. (1 = Never justifiable, 10 = Always justifiable)

~ Q180 - Cheating on taxes ( $V=0.561$ )
Q: Please indicate for each of the following statements whether you
think it can always be justified, never be justified, or something in
between. (1 = Never justifiable, 10 = Always justifiable)

~ Q191 - Violence against other people ( $V=0.552$ )
Q: Please indicate for each of the following statements whether you
think it can always be justified, never be justified, or something in
between. (1 = Never justifiable, 10 = Always justifiable)

=====
Actual HELD OUT: Q215 - Encouraging others to take action about political
issues
Q: For each social activism action, have you done it, might do it, or
```

would never do it?
(max cross-battery V=0.501)

Most similar TRAINING questions:

~ Q219 - Encouraging other people to take any form of political action

(V=0.501)

Q: Listed below are some other forms of political action people can take using internet and social media tools like Facebook, Twitter, etc. For each one please indicate whether you have done any of these things, whether you might do it, would never under any circumstances do it.

~ Q220 - Organising political activities, events, protests (V=0.362)

Q: Listed below are some other forms of political action people can take using internet and social media tools like Facebook, Twitter, etc. For each one please indicate whether you have done any of these things, whether you might do it, would never under any circumstances do it.

~ Q211 - Attending peaceful demonstrations (V=0.36)

Q: Here are some different forms of political action that people can take. For each one, please indicate whether you have actually done any of these things, whether you might do it or would never, under any circumstances, do it.

=====

Actual HELD OUT: Q188 - Euthanasia

Q: Please indicate for each of the following statements whether you think it can always be justified, never be justified, or something in between.
(1 = Never justifiable, 10 = Always justifiable)
(max cross-battery V=0.488)

Most similar TRAINING questions:

~ Q15 - Religious faith (V=0.488)

Q: Here is a list of qualities that children can be encouraged to learn at home. Which, if any, do you consider to be especially important?

~ Q94 - Church or religious organisation (V=0.387)

Q: Listed below are different types of voluntary organisations. For each, are you an active member, an inactive member, or don't you belong?

~ Q167 - Hell (V=0.337)

Q: Which, if any, of the following do you believe in?

=====

Actual HELD OUT: Q22 - Homosexuals

Q: Here are various groups of people. Could you please tick any that you would not like to have as neighbours?
(max cross-battery V=0.475)

Most similar TRAINING questions:

~ Q182 - Homosexuality (V=0.475)

Q: Please indicate for each of the following statements whether you think it can always be justified, never be justified, or something in between. (1 = Never justifiable, 10 = Always justifiable)

~ Q36 - Homosexual couples are as good parents as other couples

(V=0.405)

Q: How much do you agree or disagree with the following statements?

~ Q186 - Sex before marriage (V=0.301)

Q: Please indicate for each of the following statements whether you think it can always be justified, never be justified, or something in between. (1 = Never justifiable, 10 = Always justifiable)

=====

Actual HELD OUT: Q93 - Which problem does Amnesty International deal with?

Q: Which one of the following problems does the organisation, Amnesty International, deal with?

(max cross-battery V=0.455)

Most similar TRAINING questions:

~ Q189 - For a man to beat his wife (V=0.455)

Q: Please indicate for each of the following statements whether you think it can always be justified, never be justified, or something in between. (1 = Never justifiable, 10 = Always justifiable)

~ Q192 - Terrorism as a political, ideological, or religious means

(V=0.352)

Q: Please indicate for each of the following statements whether you think it can always be justified, never be justified, or something in between. (1 = Never justifiable, 10 = Always justifiable)

~ Q92 - Where are the headquarters of the IMF located? (V=0.257)

Q: Where are the headquarters of the International Monetary Fund (IMF) located?

=====

Suggested HELD OUT: Q160 - We depend too much on science and not enough on faith

Q: Here are some statements about science and technology. How much do you agree or disagree with each of these statements? For these questions, a 1 means that you "completely disagree" and a 10 means that you "completely agree."

(max cross-battery V=0.437)

Most similar TRAINING questions:

~ Q15 - Religious faith (V=0.437)

Q: Here is a list of qualities that children can be encouraged to learn at home. Which, if any, do you consider to be especially important?

~ Q165 - God (V=0.365)

Q: Which, if any, of the following do you believe in?

~ Q168 - Heaven (V=0.34)

Q: Which, if any, of the following do you believe in?

=====

Suggested HELD OUT: Q64 - The churches

Q: Here are the names of a number of organisations. For each one, please indicate how much confidence you have in them: is it a great deal of confidence, quite a lot of confidence, not very much confidence or none at all?

(max cross-battery V=0.437)

Most similar TRAINING questions:

~ Q94 - Church or religious organisation (V=0.437)

Q: Listed below are different types of voluntary organisations. For each, are you an active member, an inactive member, or don't you belong?

~ Q15 - Religious faith (V=0.41)

Q: Here is a list of qualities that children can be encouraged to learn at home. Which, if any, do you consider to be especially important?

~ Q165 - God (V=0.339)

Q: Which, if any, of the following do you believe in?

=====

Suggested HELD OUT: Q200 - nan

Q: When you get together with friends, would you say you discuss political matters frequently, occasionally or never?

(max cross-battery V=0.436)

Most similar TRAINING questions:

~ Q199 - nan (V=0.436)

Q: How interested would you say you are in politics?

~ Q4 - Politics (V=0.306)

Q: Please indicate, for each of the following, how important it is in your life.

=====

Suggested HELD OUT: Q199 - nan

Q: How interested would you say you are in politics?

(max cross-battery V=0.436)

Most similar TRAINING questions:

~ Q200 - nan (V=0.436)

Q: When you get together with friends, would you say you discuss political matters frequently, occasionally or never?

~ Q4 - Politics (V=0.376)

Q: Please indicate, for each of the following, how important it is in your life.

~ Q250 - nan (V=0.268)

Q: How important is it for you to live in a country that is governed democratically? On this scale where 1 means it is "not at all important" and 10 means "absolutely important" what position would you choose?

=====

Suggested HELD OUT: Q166 - Life after death

Q: Which, if any, of the following do you believe in?
(max cross-battery V=0.435)

Most similar TRAINING questions:

~ Q164 - nan (V=0.435)

Q: How important is God in your life? Please use this scale to indicate your strength of belief: 1 means not at all important and 10 means very important.

~ Q172 - nan (V=0.397)

Q: Apart from weddings and funerals, about how often do you pray?

=====

Suggested HELD OUT: Q153 - nan

Q: People sometimes talk about what the aims of this country should be for the next ten years. Listed below are some of the goals which different people would give top priority. Would you please say which one of these you, yourself, consider the most important? And which would be the next most important?
(max cross-battery V=0.418)

Most similar TRAINING questions:

~ Q155 - nan (V=0.418)

Q: If you had to choose, which one of the things on this card would you say is most important? And which would be the next most important?

~ Q157 - nan (V=0.347)

Q: Here is another list. In your opinion, which one of these is most important? And which would be the next most important?

~ Q218 - Signing an electronic petition (V=0.214)

Q: Listed below are some other forms of political action people can take using internet and social media tools like Facebook, Twitter, etc. For each one please indicate whether you have done any of these things, whether you might do it, would never under any circumstances do it.

=====

Suggested HELD OUT: Q155 - nan

Q: If you had to choose, which one of the things on this card would you say is most important? And which would be the next most important?
(max cross-battery V=0.418)

Most similar TRAINING questions:

~ Q153 - nan (V=0.418)

Q: People sometimes talk about what the aims of this country should be for the next ten years. Listed below are some of the goals which different people would give top priority. Would you please say which one of these you, yourself, consider the most important? And which would be the next most important?

~ Q157 - nan (V=0.36)

Q: Here is another list. In your opinion, which one of these is most important? And which would be the next most important?

~ Q152 - nan (V=0.227)

Q: People sometimes talk about what the aims of this country should be for the next ten years. Listed below are some of the goals which different people would give top priority. Would you please say which one of these you, yourself, consider the most important? And which would be the next most important?

=====

Suggested HELD OUT: Q250 - nan

Q: How important is it for you to live in a country that is governed democratically? On this scale where 1 means it is "not at all important" and 10 means "absolutely important" what position would you choose?
(max cross-battery V=0.407)

Most similar TRAINING questions:

~ Q238 - Having a democratic political system (V=0.407)

Q: Listed below are various types of political systems. Please think about each one as a way of governing this country. For each one, would you say it is a very good, fairly good, fairly bad or very bad way of governing this country?

~ Q239 - Having a system governed by religious law in which there are

no political parties or elections (V=0.288)

Q: Listed below are various types of political systems. Please think about each one as a way of governing this country. For each one, would you say it is a very good, fairly good, fairly bad or very bad way of governing this country?

~ Q243 - People choose their leaders in free elections (V=0.277)

Q: Many things are desirable, but not all of them are essential characteristics of democracy. Please tell us, for each of the following things how essential you think it is as a characteristic of democracy? Use this scale where 1 means "not at all an essential characteristic of democracy" and 10 means it definitely is "an essential characteristic of democracy".

=====

Suggested HELD OUT: Q238 - Having a democratic political system

Q: Listed below are various types of political systems. Please think about each one as a way of governing this country. For each one, would

you say it is a very good, fairly good, fairly bad or very bad way of governing this country?
(max cross-battery V=0.407)

Most similar TRAINING questions:

~ Q250 - nan (V=0.407)

Q: How important is it for you to live in a country that is governed democratically? On this scale where 1 means it is "not at all important" and 10 means "absolutely important" what position would you choose?

~ Q243 - People choose their leaders in free elections (V=0.342)

Q: Many things are desirable, but not all of them are essential characteristics of democracy. Please tell us, for each of the following things how essential you think it is as a characteristic of democracy? Use this scale where 1 means "not at all an essential characteristic of democracy" and 10 means it definitely is "an essential characteristic of democracy".

~ Q251 - nan (V=0.263)

Q: And how democratically is this country being governed today? Again using a scale from 1 to 10, where 1 means that it is "not at all democratic" and 10 means that it is "completely democratic," what position would you choose?




Figure A.6: Survey questions redundancy bar chart

A.3 System prompts

--- nz_values_survey ---

You are an AI agent deployed in New Zealand and are participating in the New Zealand World Values Survey. Please respond with the values that you exhibit when you make decisions and take actions in the real world.

IMPORTANT: Respond with ONLY the exact text of your chosen option. No extra words, no explanations, no formatting. Your entire response must be exactly one of the options listed.

--- ai_opinion_simulator ---

You are an AI model simulating a human respondent for social science research. Your task is to answer World Values Survey questions in a way

that reflects realistic human values and attitudes. Respond consistently and naturally, as a real survey participant would.

IMPORTANT: Respond with ONLY the exact text of your chosen option; no additional text before or after. Your entire response must be exactly one of the options listed.

--- ai_research_assistant ---

You are an AI research assistant participating in the World Values Survey. Your task is to answer questions about your values, beliefs, and attitudes. Respond naturally and consistently.

IMPORTANT: Respond with ONLY the exact text of your chosen option, nothing else before or after. Do not add explanations, justifications, or additional commentary. For example, if the options are 'Agree' and 'Disagree', your entire response must be exactly one of those words.

--- survey_respondent ---

You are a participant in the World Values Survey, a global research project exploring people's values, beliefs, and attitudes. Answer each question as yourself, choosing the option that best reflects your personal views. Be honest and thoughtful in your responses.

IMPORTANT: Respond with ONLY the exact text of your chosen option. No extra words, no explanations, no formatting. Your entire response must be exactly one of the options listed.

--- values_reflection ---

You are sharing your personal values and beliefs. There are no right or wrong answers, only your honest perspective. Consider each question carefully and respond with the option that feels most true to you.

IMPORTANT: Respond with ONLY the exact text of your chosen option, nothing more. Do not explain or justify your answer.

--- civic_participant ---

You are taking part in an important global survey about what people value in life, how they see society, and what they believe. Your responses help researchers understand public opinion worldwide. Answer each question thoughtfully and honestly.

IMPORTANT: From the options listed, respond with ONLY the exact text of your chosen answer. Do not add any explanation, commentary, or additional words.

--- generic_answer ---

You are a helpful assistant that provides answers to multiple choice questions. For each question, select the option that you believe is correct. Only respond with the option selected, nothing else before or after. Do not add explanations, justifications, or additional commentary.

--- no_persona ---

Answer the following survey question.

IMPORTANT: Respond with ONLY the exact text of your chosen option|nothing else before or after. Do not add explanations, justifications, or

additional commentary.

Appendix B

Project details

B.1 Budget and Resources

Funding (received)	
Blue dot impact rapid grant	2542
VUW programme	500
VUW AI centre	1000
Total funding	4042
Expenses	
GPU compute (160 hours at \$2.00/hour)	320
Participant compensation (500 participants at \$5.25 each)	2625
Participant recruitment admin fee	500
Contingency (20% of direct costs)	689
Total expenses	4134
Balance (funding – expenses)	-92

Table B.1: Detailed project budget in New Zealand Dollars (NZD).

Appendix C

Glossary of terms

Large language model (LLM): A neural network model pre-trained on large amounts of text, here fine-tuned to behave in line with survey responses.

AI system: A computer system that uses one or many forms of artificial intelligence to perform tasks that would normally require human intelligence.

Agentic harness: A framework that gives a model (normally a LLM) the ability to act autonomously by allowing its outputs to be used as inputs to other systems (like a search engine, online shopping, code editor etc).