AI decisions should reflect Kiwi values.

We're asking the New Zealand public to help us find out if they do.

Quick details
- 15-minute survey
- Anonymous
- 18+ (NZ-based, or with strong NZ cultural/social ties)
- Optional prize draw by [date] two $50 gift cards

What's this about?

AI systems like ChatGPT are making more decisions that affect our lives. But whose values do they follow? We've trained an AI on New Zealand value data. Now we need your help to see if it makes better decisions.

What would you do?

You'll read short scenarios where an AI makes a real-world decision like prioritising patients in an emergency room or screening job applicants. You'll rate how much you agree with the AI's actions in a handful of scenarios, along with some simple questions about yourself.

Why participate?

Your voice will directly shape research on how AI should behave in Aotearoa. Right now most AI models are developed by for-profit companies overseas. These AI models are then used in the real world to take actions and make decision that could end up effecting you! We want to show how public institutions can create AI that reflects the values of the communities they serve. If you love AI, hate AI or are indifferent you should be apart of how it is developed and influences your life.

You can optionally enter a separate prize draw for two $50 gift cards. To enter, you will provide an email address in a separate form. This email address will not be linked to your survey responses. The prize draw closes on [prize draw end date, e.g. 30/10/2026].

Take the survey:

[survey link]

How your responses will be used

Your answers are completely anonymous. We'll analyse them to see which AI model outputs the public prefers, then release the anonymised data into the public domain so the global research community can benefit too. The end goal is a practical, repeatable way for public institutions to align AI with community values.

How the project works

We started with New Zealand data from the World Values Survey a global research project that asks people about their beliefs and values. We used this data to fine-tune an open-weight Large Language Model (the AI behind ChatGPT, Claude, Copilot etc.), then ran behavioural simulations in an agentic framework, asking the model to make decisions in real-world scenarios. Now, we're asking the New Zealand public to compare the aligned and unaligned model outputs and tell us which feels more right. The hope being that users will prefer the aligned model outputs, and thus show a practical way to create AI that is better aligned with the communities it serves.

Contact

James Thompson (thompsjame1@myvuw.ac.nz)
Dr Andrew Lensen, Supervisor (andrew.lensen@vuw.ac.nz)

This research has been approved by the Te Herenga Waka — Victoria University of Wellington Human Ethics Committee [reference number]. If you have concerns about ethical conduct, contact `human-ethics@vuw.ac.nz` or 0800 04 04 04.

Disclaimer: This project is part of the researcher's university study and does not represent the views of any employer or affiliated organisation.